
BIG-GEM

The Big-GEM Theory

Heterogeneous GPU-Accelerated RTL Simulation

Native DSP48E2, CARRY4, and SRLC32E execution in NVIDIA GEM

Technical Project Report

Takneek PS – Zenith

Deliverable E

Repository: <https://github.com/Pratham-015/GEM>

Measured revision: 23fc43fac59f2ba6e349132e632cbcb062a1f39b

Report date: 2 September 2026

Project: Big-GEM / The Big-GEM Theory
Base system: NVIDIA GEM GPU-accelerated RTL simulator
Extension: Heterogeneous native execution of DSP48E2, CARRY4, and SRLC32E
Primary measured platform: NVIDIA GeForce RTX 4050 Laptop GPU (SM 8.9)
Author recorded in the measured revision: Pratham Sharma

Contents

1	Project Overview	3
1.1	Abstract	3
1.2	Problem and Contribution	3
2	Requirements and Constraints	4
3	System Architecture	4
3.1	End-to-End Data Path	4
3.2	Frontend and Macro Preservation	5
4	Heterogeneous Graph and Scheduling Equations	5
4.1	Typed Graph	5
4.2	Combinational Levelization	5
4.3	Partition and Block Assignment	6
4.4	Per-Cycle Heterogeneous Schedule	6
5	GPU Memory Hierarchy and Data Movement	7
5.1	Layout Equations	7
5.2	FPGA Primitive to GPU Hierarchy	7
6	Native Hardware Macro Semantics	8
6.1	DSP48E2 Simplified Subset	8
6.2	CARRY4 and Fused Chains	8
6.3	SRLC32E	8
7	Verification and Validation	9
8	Benchmarking and Numerical Analysis	9
8.1	Metric and Timing Boundary	9
8.2	Production Workload Throughput	10
8.3	Equivalent Upstream Comparison	10
8.4	Useful Multi-Partition Scaling	11

8.5	Nsight Compute Results	11
8.6	Unmodified GEM Comparison	11
9	Reproducibility	12
9.1	Commands	12
10	Risks, Limitations, and Failure Modes	13
11	Results Summary and PS Coverage	13
12	Conclusion	14
13	References and Project Sources	14
A	Appendix A – Exact Evidence Files	14
B	Appendix B – Final Technical Consistency Matrix	15

1. Project Overview

1.1 Abstract

Big-GEM extends NVIDIA GEM from a homogeneous one-bit And-Inverter Graph (AIG) simulator into a heterogeneous simulator that preserves and executes three FPGA primitives: DSP48E2, CARRY4, and SRLC32E. Yosys 0.68 and its Slang frontend elaborate SystemVerilog while black-box interception and a DSP-specific techmap retain macro identity. The host constructs explicit named-port macro records, classifies same-cycle and clock-boundary dependencies, fuses legal CARRY4 chains, and creates warp-padded 64-bit structure-of-arrays buffers. A custom cooperative CUDA kernel interleaves original Boomerang Boolean propagation with typed macro gather, shared-memory staging, native 64-bit evaluation, output publication, and a single synchronized state-commit boundary.

Correctness was tested through the production RTL-to-CUDA path and against independent Python, literal SystemVerilog, and installed Xilinx UNISIM models. CARRY4 covered its complete 1,024-input truth table plus randomized cases; GPU differential totals were 3,024 CARRY4 slices, 2,009 fused 60-bit chains, 2,336 DSP cycles, and 2,100 SRLC32E cycles. Four randomized production DAG topologies and the exact DSP→CARRY4→SRLC32E→DSP chain matched independent RTL with zero reported mismatches.

Production throughput was measured with 12 warm-ups and seven retained launches. The only identical-workload comparison supported by both official upstream and the modified simulator is Boolean-only: 59,232 versus 59,322 simulated cycles/s, or $1.0015 \times (+0.153\%)$. A separate 9-partition heterogeneous stress design scaled from 5,273 cycles/s at one block to 16,677 cycles/s at 12 blocks ($3.163 \times$). No valid heterogeneous upstream speedup is claimed: the historical shredded-netlist comparison failed its RTL correctness gate.

1.2 Problem and Contribution

The problem statement identifies a structural inefficiency in conventional GEM: Yosys decomposition converts wide arithmetic and stateful primitives into large one-bit networks. Big-GEM changes the representation and execution path rather than adding an isolated macro microbenchmark. Its concrete contributions are:

1. a Yosys 0.68 macro-preserving and DSP-normalizing frontend;
2. a heterogeneous host graph with stable port identities, dependency timing, and explicit combinational levels;
3. separate bit-packed Boolean storage and 64-bit aligned, warp-padded macro state/I/O storage;
4. a custom CUDA schedule that places macro operations between Boolean settling barriers and commits DSP/SRL state at one global edge; and
5. reproducible differential, throughput, block-scaling, and Nsight Compute evidence.

Evidence boundary. Results in this report were regenerated from commit 23fc43f while the working tree contained the scheduler/benchmark changes listed by Git as uncommitted. Every table therefore records the revision and dirty-tree status. The repository should be committed and the benchmark command rerun before final archival publication.

2. Requirements and Constraints

Table 1: Problem-statement requirements and demonstrated status.

Requirement	Status	Evidence in the current repository
Preserve DSP48E2, CARRY4, SRLC32E through Yosys	Verified	Production JSON checks in <code>verif/full_integration_test.py</code> ; Yosys 0.68 observed.
Mixed 1-bit and 48-bit heterogeneous scheduling	Verified	Typed dependency graph and production exact-chain/random-DAG differentials.
64-bit aligned contiguous macro buffers	Verified	<code>src/macro_layout.rs</code> ; four layout tests, including warp-scale non-overlap.
Native GPU ALU macro evaluation	Verified	<code>csrc/gem_macros.cuh</code> executed on GPU and differentially checked.
Correct macro-to-macro sequencing	Verified	Explicit levels; fused CARRY4 chains; exact-chain and randomized production tests.
Shared memory and warp synchronization	Verified	12,288-byte macro tile, <code>__ballot_sync</code> , block barriers, and cooperative grid barriers.
Yosys 0.68 and IEEE 1800-2012 syntax	Verified	Installed Yosys 0.68; Slang parses the SystemVerilog-2012 feature test.
Single global clock and zero initialization	Verified	Multiple macro clocks are rejected; state buffers are zero-created and checked.
Throughput versus unmodified GEM	Partial	Identical Boolean workload measured; no valid equivalent heterogeneous upstream comparison.
Nsight bandwidth/divergence report	Verified	Four current production-kernel profiles are stored as raw CSV and parsed JSON.

The required DSP configuration is enforced structurally: AREG, BREG, CREG, DREG, ADREG, and MREG are zero and PREG is one. Unsupported parameterizations cause the DSP techmap to fail instead of being silently simulated with different latency. The simulator assumes one global clock domain and initializes all internal macro state to zero; INIT-string parsing remains intentionally out of scope, exactly as permitted by the problem statement.

3. System Architecture

3.1 End-to-End Data Path




Figure 1: Actual Big-GEM production path. No standalone macro microkernel is substituted for the measured simulator.

3.2 Frontend and Macro Preservation

`scripts/synthesize_macros.py` checks the installed version string, invokes `read_slang` with an IEEE 1800-2017 mode that accepts the required 2012 syntax, elaborates the requested top, and runs synthesis/ABC around black-box macro declarations. `aigpdk/dsp48e2_ps_map.v` transforms a supported raw DSP48E2 into `GEM_DSP48E2` with two state bits and one pre-adder-control bit. `CARRY4` and `SRLC32E` remain named black boxes. JSON is not the runtime parser input; it is retained as machine-checkable preservation evidence. The runtime host parser consumes the emitted structural Verilog and constructs `NetlistDB`, `AIG`, and macro records.

The DSP control mapping is:

$$\text{state(OPMODE)} = \begin{cases} 0, & \text{OPMODE} = 0\text{x}030, \\ 1, & \text{OPMODE} = 0\text{x}005, \\ 2, & \text{OPMODE} = 0\text{x}025, \end{cases} \quad \text{use_pre} = \text{INMODE}[2].$$

Plain-add `ALUMODE` and supported `INMODE` encodings are recognized by the map. Static pipeline parameters outside the required subset are rejected.

4. Heterogeneous Graph and Scheduling Equations

4.1 Typed Graph

Let the execution graph be

$$G = (V, E), \quad V = V_B \cup V_C \cup V_S \cup V_D,$$

where V_B contains one-bit AIG operations, V_C fused `CARRY4`-chain nodes, V_S `SRLC32E` nodes, and V_D `DSP48E2` nodes. The sets are disjoint and each macro instance retains named input/output ports and bit indices. Define

$$T(v) \in \{\text{AIG}, \text{CARRY}, \text{SRL}, \text{DSP}\}.$$

An edge records producer cell/port/bit and consumer cell/port/bit. The frontend traces through intervening AIG logic, so a macro dependency is recovered even when ordinary Boolean nodes lie between producer and consumer.

Each macro input has role $r_i \in \{\text{comb}, \text{next}\}$ and each macro output has role $r_o \in \{\text{comb}, \text{reg}\}$. The implemented timing classifier is exactly

$$\tau(e) = \begin{cases} \text{same}, & r_i(e) = \text{comb} \wedge r_o(e) = \text{comb}, \\ \text{edge}, & \text{otherwise.} \end{cases}$$

`CARRY O/CO` and `SRL Q/Q31` are combinational outputs; `DSP P` is registered. `CARRY` inputs and `SRL` address are combinational; `DSP` operands/control and `SRL D/CE` are next-state inputs.

4.2 Combinational Levelization

Only nodes with combinational outputs participate in the same-cycle macro DAG $G_c = (V_c, E_c)$, where $E_c = \{e \in E : \tau(e) = \text{same}\}$. The host uses Kahn levelization. Its equivalent recurrence is

$$\ell(v) = \begin{cases} 0, & \text{pred}_{E_c}(v) = \emptyset, \\ 1 + \max_{u \in \text{pred}_{E_c}(v)} \ell(u), & \text{otherwise.} \end{cases}$$

The cohort at level k is $M_k = \{v \in V_c : \ell(v) = k\}$. A nonempty residual indegree after Kahn traversal is rejected as a combinational macro cycle. DSP nodes do not enter G_c because PREG=1 makes P a cycle-boundary source.

Legal adjacent CARRY4 instances are fused before levelization when CO[3] drives the next CIN, CYINIT is inactive, and total width is at most 60 bits. For a fused chain, the internal macro-to-macro edges disappear and one native operation publishes all O/CO bits. Non-fused same-cycle macro edges are ordered by ℓ and communicated through explicitly allocated output-bit slots in the packed Boolean-state row. Thus ordering is explicit, although the current implementation does not forward a producer's 64-bit register value directly to a consumer register.

4.3 Partition and Block Assignment

For endpoint cone i , let w_i be the host-computed fractional AIG work rounded to an integer. When the optional external hypergraph partitioner is absent, the current fallback applies deterministic longest-processing-time assignment:

$$i_1, \dots, i_n = \text{argsort}_i(-w_i, i), \quad p(i_j) = \arg \min_{q \in [0, k)} (W_q, q),$$

$$W_{p(i_j)} \leftarrow W_{p(i_j)} + \max(w_{i_j}, 1).$$

Mapping validity remains controlled by `Partition::build\_one`; failing groups are repartitioned until each is legal. Effective partitions are assigned to CUDA blocks with a second greedy load rule. If $d(p)$ is the number of Boomerang stages in partition p , its placement cost is $c(p) = d(p) + 3$ and

$$b(p) = \arg \min_b L_b, \quad L_{b(p)} \leftarrow L_{b(p)} + c(p).$$

4.4 Per-Cycle Heterogeneous Schedule

Let $\mathcal{B}(x_t, y)$ denote one complete pass over all major Boolean Boomerang stages using cycle inputs x_t and mutable output row y . Let $\mathcal{G}_k$, $\mathcal{E}_k$, and $\mathcal{W}_k$ be macro gather, evaluate, and publish operations for level k . For state $q_t = (P_t, R_t)$:

1. publish old DSP registered outputs P_t ;
2. for phase $h = 0, 1$, and every $k = 0, \dots, K - 1$, execute

$$y \leftarrow \mathcal{B}(x_t, y), \quad z_k \leftarrow \mathcal{G}_k(y), \quad m_k \leftarrow \mathcal{E}_k(z_k, q), \quad y \leftarrow \mathcal{W}_k(m_k, y),$$

followed by one final $y \leftarrow \mathcal{B}(x_t, y)$;

3. between phases ($h = 0$ only), gather all DSP and SRL next-state inputs from the same settled snapshot and commit the single global edge:

$$q_{t+1} = F(q_t, x_t, y^-);$$

4. phase $h = 1$ propagates newly registered P and SRL Q/Q31 values to same-cycle observable outputs and downstream next-state inputs.

Every gather/evaluate/publish boundary is separated by a cooperative grid barrier. This is a deterministic two-phase level walk, not an unconstrained Jacobi loop. Its cost limitation is also explicit: the Boolean network is re-executed $2(K + 1)$ times per simulated cycle for heterogeneous designs.

5. GPU Memory Hierarchy and Data Movement

5.1 Layout Equations

Boolean state remains bit-packed in 32-bit words. Macro state and I/O occupy separate CUDA allocations of 64-bit words. For kind t with n_t instances, the padded stride is

$$s_t = 32 \left\lceil \frac{n_t}{32} \right\rceil.$$

The structure-of-arrays address of field f for instance i is

$$\text{addr}(t, f, i) = B_t + 8(f s_t + i) \quad \text{bytes.}$$

The allocator uses eight I/O fields for CARRY and DSP and four for SRL. DSP P and the SRL's 32-bit state each consume one warp-padded 64-bit state slot; CARRY is stateless. CUDA allocations and every field address are therefore 8-byte aligned, and adjacent lanes access adjacent 64-bit words for one field.

5.2 FPGA Primitive to GPU Hierarchy




Figure 2: Actual macro data movement. Global memory is the grid-wide exchange boundary; shared memory is a block-local staging tile; arithmetic temporaries remain in registers.

Table 2: Primitive-specific mapping onto the measured CUDA hierarchy.

Primitive	Datapath	State	Global memory	Shared tile	Registers / execution
DSP48E2	27+27, 18, 45, 48 bit	48-bit P	Five SoA input fields, P state, output bit slots	Five inputs + one state word/thread	Signed 64-bit pre-add, multiply, and masked ALU select
CARRY chain	4–60 bit	None	Three input fields and O/CO output slots	S, DI, CIN words/thread	One 64-bit add plus Boolean transforms
SRLC32E	5-bit address, 1-bit I/O	32-bit SRL	One packed control field, state word, Q/Q31 slots	Control + state word/thread	Shift/mask and dynamic register read

The full measured kernel uses 16,640 bytes of static shared memory per block: 12,288 bytes are the dedicated macro tile; the remainder is Boolean metadata, writeout, state, and script-index

storage. Nsight reports 124 registers/thread. The kernel uses `__ballot\__sync` to form selected macro cohorts and `__all\__sync` to assert type uniformity. Block-local phases use `__syncthreads;` grid-wide visibility uses `cooperative\groups::this\grid().sync()`.

6. Native Hardware Macro Semantics

6.1 DSP48E2 Simplified Subset

All operands are interpreted as signed two's-complement values after masking to their declared widths. With sext_w denoting sign extension and $[x]_w = x \bmod 2^w$:

$$AD = \text{sext}_{27}([A + uD]_{27}), \quad M = AD \cdot \text{sext}_{18}(B),$$

where $u \in \{0, 1\}$ is `use\_pre`. At the single rising edge,

$$P_{t+1} = \left[\begin{cases} \text{sext}_{48}(C), & s = 0, \\ M, & s = 1, \\ \text{sext}_{48}(P_t) + M, & s \geq 2 \end{cases} \right]_{48}.$$

The CUDA implementation computes this with native signed 64-bit arithmetic and branch-free all-zero/all-one select masks. Overflow therefore wraps exactly at 27 bits in the pre-adder and 48 bits in P.

6.2 CARRY4 and Fused Chains

The literal specification is

$$C_0 = \text{CYINIT} \vee \text{CIN}, \quad C_{i+1} = (S_i \wedge C_i) \vee (\neg S_i \wedge DI_i),$$

$$O_i = S_i \oplus C_i, \quad CO_i = C_{i+1}.$$

For a fused n -bit chain, the implementation defines

$$A' = S \vee DI, \quad B' = \neg S \wedge DI, \quad T = A' + B' + C_0,$$

$$O = [T]_n, \quad CO = [(T \oplus A' \oplus B') \gg 1]_n.$$

This identity was checked against the literal ripple equations. The maximum width is 60 bits so the final carry remains representable in the 64-bit sum.

6.3 SRLC32E

For state $R_t \in \{0, 1\}^{32}$, address a , data d , and clock enable e :

$$R_{t+1} = \begin{cases} [(R_t \ll 1) \vee d]_{32}, & e = 1, \\ R_t, & e = 0, \end{cases} \quad Q = R[a \bmod 32], \quad Q_{31} = R[31].$$

DSP and SRL next-state inputs are gathered before either state array is updated, which gives both primitives the same logical rising edge. Persistent state is allocated with zero initialization.

7. Verification and Validation

Table 3: Executed verification evidence on 2 September 2026.

Test	Result	Coverage
Rust workspace, all targets, CUDA feature	PASS	1 partitioner unit test; 4 memory tests; 4 parser/DAG integration tests; all binaries compiled.
CARRY4 differential	PASS	3,024 slices, including complete 1,024-case truth table; 2,009 60-bit chains.
DSP48E2 differential	PASS	2,336 sequential vectors with signed/boundary/random cases.
SRLC32E differential	PASS	2,100 sequential vectors with CE/address changes.
Independent references	PASS	Python, literal SystemVerilog/Icarus, installed Xilinx UNISIM, and production CUDA.
Integrated RTL pipeline	PASS	RTL→Yosys→JSON→partition→CUDA→VCD.
Exact required chain	PASS	DSP→CARRY4→SRLC32E→DSP on 1 and 4 blocks; 2,709 values checked/run, zero mismatches.
Random production DAGs	PASS	4 seeded topologies, 192 stimulus cycles, 1,029 values/run, zero mismatches.
Wide multi-partition path	PASS	9 partitions, 16 blocks, 18 complete 8,192-bit result frames, zero mismatches.
Nsight production profiles	PASS	Exactly one production simulation kernel selected per profile.
Compute Sanitizer	BLOCKED	WDDM debugger initialization failed and this device/-driver combination was reported unsupported.

The production simulator’s optional `-check-with-cpu` path provides a second schedule/state sanity check. Expected external outputs are nevertheless generated independently by Verilog/UNISIM rather than by reusing the CUDA evaluator. The full regression command completed successfully:

```
python3 verif/full_integration_test.py
```

8. Benchmarking and Numerical Analysis

8.1 Metric and Timing Boundary

The primary metric is

$$\text{CPS} = \frac{N_{\text{simulated cycles}}}{t_{\text{launch+kernel+sync}}}.$$

Each workload is synthesized and partitioned once. The measured interval wraps one production cooperative-kernel launch and mandatory device synchronization; it excludes parsing, allocation, H→D/D→H transfer, synthesis, partitioning, and output generation. Twelve launches are discarded before seven measurements. State and deterministic inputs are reset outside each measured interval. The table uses the median CPS; “runtime” is calculated as cycles/median CPS, while CV is sample standard deviation divided by mean CPS.

8.2 Production Workload Throughput

Table 4: Current production-kernel throughput (seven retained samples).

Workload	Cycles	Blocks	AIG	Macros	D/C/S	Runtime (ms)	Median CPS
Boolean heavy	2,000	4	507		0/0/0	25.445	78,600
DSP heavy	2,000	4	4,464		32/0/0	95.050	21,041
Carry heavy	1,000	4	2,880		0/128/0	636.486	1,571
SRL heavy	2,000	4	576		0/0/128	46.847	42,692
Heterogeneous small	1,000	4	1,120		4/8/8	129.535	7,720
Heterogeneous mixed	1,000	4	3,520		16/32/32	281.028	3,558
Deep dependency	200	4	1,981		0/64/0	428.885	466
Heterogeneous large	200	4	8,704		32/128/128	190.018	1,053
Occupancy stress	1,000	16	82,385		1/1/1	60.004	16,666

The data show workload cost, not a cross-workload speedup. Carry-heavy and deep dependency cases are slow because every combinational macro level forces another complete Boolean settling pass and grid barriers. SRL-heavy is much faster than carry-heavy despite equal macro count because independent SRLs share a shallow level. Mixed scaling drops from 7,720 CPS (small) through 3,558 CPS (medium) to 1,053 CPS (large) as AIG and macro work increase. These values demonstrate functional scaling but do not isolate the benefit of preserving a macro.

All seven-sample coefficients of variation were below 0.43%; mixed and large were 0.015% and 0.030%, respectively. This repeatability does not remove systematic issues such as laptop power state or the dirty-tree provenance.

8.3 Equivalent Upstream Comparison

Table 5: Identical macro-free Boolean workload on official upstream and modified GEM.

Implementation	Cycles	Mean runtime (ms)	Median CPS	CPS CV
Official upstream 9e913f9	2,000	33.779	59,231.5	0.361%
Modified Big-GEM 23fc43f	2,000	33.712	59,322.3	0.132%

Measured and calculated values are separated as follows:

$$\text{speedup} = \frac{59,322.260}{59,231.513} = 1.001532\times, \quad \text{gain} = 0.1532\%.$$

This is parity within a small fraction of one percent, not evidence of a macro speedup. It does show that the Boolean-only fast path preserves upstream throughput on this workload. The benchmark runner attempted a same-RTL shredded-upstream versus preserved-modified experiment, but at least one representation failed the independent RTL differential. Its performance number was therefore suppressed. No heterogeneous baseline speedup is reported.

8.4 Useful Multi-Partition Scaling

Table 6: Nine-partition occupancy-stress block sweep, 1,000 cycles, seven samples.

Blocks	1	4	8	9	12	16	20	28/40
CPS	5,273	10,776	13,056	16,670	16,677	16,665	16,661	11,452/11,451

The best retained median occurs at 12 blocks:

$$\frac{16,676.956}{5,272.581} = 3.16296 \times \quad \text{or} \quad 216.30\% \text{ gain over one block.}$$

Only nine blocks have useful partition work. Additional cooperative blocks are idle but still participate in grid synchronization; performance is flat through 20 blocks and falls by about 31% at 28–40 blocks. This is direct evidence that partitioning exposes useful parallel work, and equally direct evidence that oversizing the cooperative grid is harmful.

8.5 Nsight Compute Results

Table 7: Current production-kernel Nsight Compute counters.

Workload	Grid	Ach./theor. occ.	Divergent targets	Predicated lanes	DRAM MB/s	LD/ST sectors
Exact chain	1	16.67/33.33%	3.764%	70.09%	6.52	6.63/2.89
Mixed	4	16.67/33.33%	1.212%	75.16%	1.95	7.36/3.73
Large	4	16.67/33.33%	1.078%	75.59%	2.73	7.52/3.70
Occupancy stress	16	16.71/33.33%	1.306%	83.22%	15.19	9.72/4.06

The profiled symbol was `simulate\_v1\_noninteractive\_simple\_scan`; each profile selected exactly one production kernel. The environment was Nsight Compute 2025.2.1.0, CUDA 12.9, driver 596.49, and RTX 4050 Laptop GPU. DRAM peak utilization rounded to 0.00–0.01%, so these tests are not bandwidth-saturated. The principal limits are synchronization, insufficient useful blocks in three workloads, and a 124-register/thread kernel whose theoretical occupancy is 33.33%. Branch-target divergence is low (1.08–3.76%), but predicated-lane utilization of 70–83% shows that branch uniformity is not equivalent to full lane efficiency.

The load/store sector ratios are transaction-density counters, not a direct coalescing-efficiency percentage because the kernel mixes access widths and instruction sites. They therefore support neither a claim of perfect coalescing nor a bandwidth bottleneck. The SoA address mapping is structurally contiguous; instruction-level coalescing attribution remains future profiling work.

8.6 Unmodified GEM Comparison

Table 8: Controlled comparison against unmodified GEM.

Pool / implementation	CPS	Speedup
Official upstream, identical Boolean workload	59,231.5	1.000×
Our submission, identical Boolean workload	59,322.3	1.0015×

9. Reproducibility

Table 9: Measured software and hardware configuration.

Item	Recorded value
Repository	https://github.com/Pratham-015/GEM
Git revision	23fc43fac59f2ba6e349132e632cbcb062a1f39b; dirty working tree
GPU / compute capability	NVIDIA GeForce RTX 4050 Laptop GPU / 8.9
Driver / VRAM	596.49 / 6,141 MiB
CUDA compiler	CUDA 12.9, NVCC 12.9.86
Nsight Compute	2025.2.1.0
Yosys	0.68, git 38e001a6f, Slang frontend
Rust	rustc 1.98.0
CUDA launch	256 threads/block; cooperative production kernel

9.1 Commands

```
python3 verif/full_integration_test.py
python3 benchmark/run_benchmarks.py --repetitions 7
python3 benchmark/profile_boomerang_ncu.py
python3 benchmark/profile_boomerang_ncu.py \
  --workload mixed_heterogeneous --skip-prepare
python3 benchmark/profile_boomerang_ncu.py \
  --workload large_scale --skip-prepare
python3 benchmark/profile_boomerang_ncu.py \
  --workload occupancy_stress --blocks 16 --skip-prepare
python3 benchmark/profile_block_sweep.py \
  --workload occupancy_stress --repetitions 7
```

Machine-readable results are in `benchmark/results/latest.json`, `benchmark/results/block_sweep_occupancy_stress.json`, and `benchmark/nsight_*.json`; raw Nsight CSV files are stored beside the JSON.

10. Risks, Limitations, and Failure Modes

Table 10: Known limitations at report generation time.

Limitation	Consequence / mitigation	Status
No valid heterogeneous upstream differential	Native-preservation speedup cannot be claimed; repair the shredded golden flow and rerun equivalent RTL.	Open
Dirty-tree measurements	Commit hash alone cannot reconstruct every measured byte; commit changes and rerun archival evidence.	Open
Two full Boolean walks per macro phase	Deep combinational macro levels incur repeated AIG work and grid barriers.	Measured
Macro exchange uses global bit slots	Correct across blocks, but direct word-level macro forwarding is not implemented for non-fused edges.	Architectural limit
124 registers/thread	Theoretical occupancy is 33.33%; small grids achieve about 16.67%.	Measured
Compute Sanitizer unavailable	Memory/race diagnostics were not obtained from this WDDM/driver/device combination.	Blocked
SystemVerilog breadth	Valid hidden designs beyond the exercised hierarchy/generate/signed tests may expose parser assumptions.	Residual risk

No result supports the phrases “zero divergence”, “fully coalesced”, or “macro acceleration versus upstream”. The report intentionally does not make those claims.

11. Results Summary and PS Coverage

Table 11: Deliverable E coverage.

Deliverable E requirement	Coverage	Location
Mathematical scheduling definition	Complete	Section 4: typed graph, edge timing, level recurrence, partition/block assignment, and two-phase cycle equation.
FPGA-to-GPU architecture diagrams	Complete	Figures 1–2 and primitive memory-mapping table.
CUDA architecture and memory layout	Complete	Sections 3–6, including global/shared/register placement and synchronization.
Extensive numerical analysis	Complete with limitation	Section 8: nine workloads, upstream parity, block scaling, Nsight, uncertainty, and missing external data.
Benchmarking under a unique header	Complete	Section 8 is titled “Benchmarking and Numerical Analysis”.
Reproducibility	Complete	Section 9 records versions, revision, commands, and evidence paths.

12. Conclusion

The current repository contains a real integrated heterogeneous execution path: SystemVerilog reaches Yosys 0.68, named macros survive synthesis, the host builds typed instances and dependencies, the allocator creates aligned SoA storage, and the production cooperative CUDA kernel executes Boolean and native macro phases with explicit synchronization and clock-state semantics. The strongest evidence is correctness: exhaustive/differential primitive tests, independent UNISIM and RTL references, exact required-chain simulation, randomized production DAGs, and a wide multi-partition simulation all pass.

Performance evidence is more limited. The extension preserves Boolean-only throughput within 0.153% of upstream and exposes a $3.163\times$ useful scaling gain on a 9-partition stress graph. Nsight demonstrates low branch-target divergence but low absolute bandwidth and occupancy constrained by workload size and register use. Because the heterogeneous shredded baseline failed correctness, the submission does not yet demonstrate a valid native-macro speedup over upstream GEM. This is the most important performance-evidence gap before judging.

13. References and Project Sources

1. Takneek PS – Zenith, *The Big-GEM Theory*, supplied six-page problem statement, 2026.
2. Z. Guo, Y. Zhang, R. Wang, Y. Lin, and H. Ren, “GEM: GPU-Accelerated Emulator-Inspired RTL Simulation,” DAC 2025.
3. NVIDIA, official GEM repository, <https://github.com/NVlabs/GEM>, upstream comparison commit 9e913f9b5efc8b12027bfb374be8b1a0028df00a.
4. Big-GEM implementation sources: `src/aig.rs`, `src/flatten.rs`, `src/macro_layout.rs`, `csrc/kernel_v1_impl.cuh`, and `csrc/gem_macros.cuh`.
5. Verification and measurement sources: `verif/full_integration_test.py`, `verif/host/diff_harness.py`, and `benchmark/run_deliverable_d.py`.

A. Appendix A – Exact Evidence Files

Evidence	Purpose
<code>benchmark/results/latest.json</code>	Raw commands, environment, seven samples/workload, summaries, and upstream comparison.
<code>benchmark/nsight_boomerang.csv/json</code>	Exact-chain production-kernel counters.
<code>benchmark/nsight_mixed_heterogeneous.csv/json</code>	Mixed production-kernel counters.
<code>benchmark/nsight_large_scale.csv/json</code>	Large production-kernel counters.
<code>benchmark/nsight_occupancy_stress.csv/json</code>	Multi-partition stress production-kernel counters.
<code>benchmark/results/block_sweep_occupancy_stress.json</code>	Per-block-count raw stdout and seven throughput samples.

Evidence	Purpose
<code>verif/host/diff_harness.py</code>	Independent CPU/RTL/UNISIM/GPU primitive differential harness.
<code>verif/full_integration_test.py</code>	Full synthesis, parser, scheduler, CUDA, VCD, randomized, and Nsight regression.

B. Appendix B – Final Technical Consistency Matrix

Report claim	Source implementation	Executed evidence
Macro preservation	<code>scripts/synthesize_macros.py</code> , <code>aigpdk/dsp48e2_ps_map.v</code>	JSON cell-count assertions
Same-cycle macro levels	<code>src/aig.rs</code>	exact-chain/random-DAG VCD differentials
64-bit SoA storage	<code>src/macro_layout.rs</code>	<code>tests/macro_memory_layout_test.rs</code>
Shared macro tile	<code>csrc/kernel_v1_impl.cuh</code>	cuobjdump resource assertion; Nsight 16,640 B/block
Global edge semantics	gather-before-commit CUDA schedule	DSP/SRL multi-cycle and integrated RTL tests
Throughput numbers	<code>benchmark/run_deliverable_d.py</code>	<code>benchmark/results/latest.json</code>
Warp/memory counters	<code>benchmark/profile_boomerang_ncu.py</code>	raw Nsight CSV + parsed JSON